



- **Namespaces:**
  - Partition resources and provide isolation within a cluster.
  - Logically divide a cluster into multiple virtual clusters.
  - Each namespace has its own set of resources.
- **Resource Quotas:**
  - Limit the amount of resources that can be used within a namespace.
  - Help ensure that a namespace does not consume excessive resources.
  - Set limits on CPU, memory, and storage usage within a namespace.
  - Specify limits on the number of pods, services, and other resources that can be created.
- **Together:**
  - Effectively manage and allocate resources within a Kubernetes cluster.
  - Provide better isolation and control over resource usage.

## Benefits of using namespaces and resource quotas

- Improve resource utilization
- Prevent resource conflicts
- Improve security
- Make it easier to manage large Kubernetes deployments

Lets do R & D ,

Create a namespace for testing,

Namespace name: vishalk17

- `kubectl create ns vishalk17`

```
vagrant@k8s-master:~$ kubectl create ns vishalk17
namespace/vishalk17 created
vagrant@k8s-master:~$ 
vagrant@k8s-master:~$ kubectl get ns
NAME                STATUS   AGE
kube-system          Active   21h
kube-public           Active   21h
kube-node-lease       Active   21h
default               Active   21h
vishalk17             Active   7s
vagrant@k8s-master:~$
```

## 1. Setting Resource Quotas for a Namespace:

Use the `ResourceQuota` YAML file (`quota.yaml`) when you want to enforce resource limits for a specific namespace. This can be useful in scenarios where you want to allocate a specific amount of resources to a namespace, preventing it from consuming more than the allocated resources. It helps in resource allocation control and ensuring fair sharing of resources among different namespaces.

Kubernetes doc: <https://kubernetes.io/docs/concepts/policy/resource-quotas/>

When several users or teams share a cluster with a fixed number of nodes, there is a concern that one team could use more than its fair share of resources.

Resource quotas are a tool for administrators to address this concern.

A resource quota, defined by a `ResourceQuota` object, provides constraints that limit aggregate resource consumption per namespace. It can limit the quantity of objects that can be created in a namespace by type, as well as the total amount of compute resources that may be consumed by resources in that namespace.

### Compute Resource Quota

You can limit the total sum of `compute resources` that can be requested in a given namespace.

The following resource types are supported:

Resource Name	Description
<code>limits.cpu</code>	Across all pods in a non-terminal state, the sum of CPU limits cannot exceed this value.
<code>limits.memory</code>	Across all pods in a non-terminal state, the sum of memory limits cannot exceed this value.
<code>requests.cpu</code>	Across all pods in a non-terminal state, the sum of CPU requests cannot exceed this value.
<code>requests.memory</code>	Across all pods in a non-terminal state, the sum of memory requests cannot exceed this value.
<code>hugepages-&lt;size&gt;</code>	Across all pods in a non-terminal state, the number of huge page requests of the specified size cannot exceed this value.
<code>cpu</code>	Same as <code>requests.cpu</code>
<code>memory</code>	Same as <code>requests.memory</code>

## **\*\* Configure Memory and CPU Quotas for a Namespace \*\***

how to set quotas for the total amount memory and CPU that can be used by all Pods running in a namespace.

### **Create a ResourceQuota**

```
apiVersion: v1
kind: ResourceQuota
metadata:
  name: mem-cpu-demo
  namespace: vishalk17
spec:
  hard:
    requests.cpu: "1"
    limits.cpu: "2"

    requests.memory: 500Mi
    limits.memory: 1Gi
```

- save , exit, apply

The **requests.cpu** and **requests.memory** fields specify the minimum amount of CPU and memory that each pod in the namespace must request. If a pod does not specify a request for a resource, the scheduler will automatically assign the minimum request for that resource.

The **limits.cpu** and **limits.memory** fields specify the maximum amount of CPU and memory that each pod in the namespace can consume. If a pod tries to consume more than its limit, the kubelet will kill the pod.

```
vagrant@k8s-master:~$ kubectl get resourcequota -n vishalk17
NAME          AGE    REQUEST                                     LIMIT
mem-cpu-demo  97s    requests.cpu: 0/1, requests.memory: 0/500Mi  limits.cpu: 0/2, limits.memory: 0/16i
vagrant@k8s-master:~$
vagrant@k8s-master:~$ kubectl describe resourcequota -n vishalk17
Name:          mem-cpu-demo
Namespace:     vishalk17
Resource       Used    Hard
-----
limits.cpu     0      2
limits.memory  0      16i
requests.cpu   0      1
requests.memory 0      500Mi
vagrant@k8s-master:~$
```

So, all has been set it seems

## Situation 1 : What if i dont set limit and request in pod manifest file

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
```

- Save , apply

```
vagrant@k8s-master:~$
vagrant@k8s-master:~$ kubectl apply -f nginx-pod.yml -n vishalk17
Error from server (Forbidden): error when creating "nginx-pod.yml": pods "nginx" is forbidden: failed quota: mem-cpu-demo: must specify
limits.cpu for: nginx; limits.memory for: nginx; requests.cpu for: nginx; requests.memory for: nginx
vagrant@k8s-master:~$
```

### Noticeable Changes :

- Pod will not create until and unless you specify limit and request in manifest file of pod

## Situation 2 : add limit and request to the pod

### 2.1 keeping limit and request value lower compared to resource quota

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      limits:
        memory: "300Gi"
        cpu: "800m"
      requests:
        memory: "200Mi"
        cpu: "200m"
```

- Save and apply

```
vagrant@k8s-master:~$ kubectl apply -f nginx-pod.yml -n vishalk17
Error from server (Forbidden): error when creating "nginx-pod.yml": pods "nginx" is forbidden: exceeded quota: mem-cpu-demo, requested:
limits.memory=300Gi, used: limits.memory=0, limited: limits.memory=16i
vagrant@k8s-master:~$
```

Noticeable Changes :

- Pod will not create I assume its because of the limit we have lower below requested the one defined in RQ

## 2.2 keeping request equal and limits equal compared to RQ

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      limits:
        memory: "1Gi"
        cpu: "2"
      requests:
        memory: "500Mi"
        cpu: "1"
```

```
vagrant@k8s-master:~$ kubectl apply -f nginx-pod.yml -n vishalk17
pod/nginx created
vagrant@k8s-master:~$ kubectl get pods -n vishalk17
NAME     READY   STATUS    RESTARTS   AGE
nginx    1/1     Running   0           8s
vagrant@k8s-master:~$ kubectl get resourcequota -n vishalk17
NAME      AGE   REQUEST              LIMIT
mem-cpu-demo 32m   requests.cpu: 1/1, requests.memory: 500Mi/500Mi  limits.cpu: 2/2, limits.memory: 1Gi/1Gi
vagrant@k8s-master:~$
```

Noticeable Changes :

- Pod created

## 2.3 keeping request equal and limits lower compared to RQ

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx
spec:
  containers:
  - name: nginx
    image: nginx
    resources:
      limits:
        memory: "800Mi"
        cpu: "1.5"
      requests:
        memory: "500Mi"
        cpu: "1"
```



- Save, apply

```
vagrant@k8s-master:~$  
vagrant@k8s-master:~$ kubectl apply -f nginx-pod.yml -n vishalk17  
pod/nginx created  
vagrant@k8s-master:~$ kubectl get pods -n vishalk17  
NAME      READY   STATUS    RESTARTS   AGE  
nginx     1/1     Running   0           5s  
vagrant@k8s-master:~$
```

Noticeable Changes :

- Pod created

Summary, in case you have defined RQ for namespace,

- In pod specification request must be equal to request defined in RQ else pod will not create
- Acceptable : In pod specification limit can be lower or equal to limit defined in RQ
- Pod will not create : if no limits defined in pod specification.

### Situation 3 : What if I dont set Resource Quota for namespace and add limit and request to the pod

Check if I have any resource quota applied to the namespace or not , if applied then I will remove to achieve this situation.

```
vagrant@k8s-master:~$  
vagrant@k8s-master:~$ kubectl get resourcequota -n vishalk17  
No resources found in vishalk17 namespace.  
vagrant@k8s-master:~$
```

Good to go.

First of all lets check how much cpu and memory using by particular container in a pod

- vi httpd.yml

```
apiVersion: apps/v1  
kind: Deployment  
metadata:  
  name: apache-deployment  
  labels:  
    app: apache  
spec:  
  replicas: 1  
  selector:  
    matchLabels:  
      app: apache  
  template:  
    metadata:  
      labels:  
        app: apache  
    spec:  
      containers:  
      - name: apache  
        image: httpd:latest  
        ports:  
        - containerPort: 80
```

- Save and apply

```
vagrant@k8s-master:~$  
vagrant@k8s-master:~$ kubectl apply -f httpd.yml -n vishalk17  
deployment.apps/apache-deployment created  
vagrant@k8s-master:~$  
vagrant@k8s-master:~$ kubectl top pods -n vishalk17  
NAME                                CPU(cores)  MEMORY(bytes)  
apache-deployment-6696dd9569-t5wkc  1m          6Mi  
vagrant@k8s-master:~$
```

## Situation 3.1 : What if I set limit memory to 4Mi and keep cpu same as 1m

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apache-deployment
  labels:
    app: apache
spec:
  replicas: 1
  selector:
    matchLabels:
      app: apache
  template:
    metadata:
      labels:
        app: apache
    spec:
      containers:
        - name: apache
          image: httpd:latest
          ports:
            - containerPort: 80
          resources:
            limits:
              memory: "4Mi"
              cpu: "10m"
            requests:
              memory: "3Mi"
              cpu: "1m"
```

```
vagrant@k8s-master:~$ kubectl get all -n vishalk17
NAME                                READY    STATUS              RESTARTS   AGE
pod/apache-deployment-6d9b8fb897-f9k4m  0/1      CrashLoopBackOff    3 (26s ago)  102s

NAME                                READY    UP-TO-DATE    AVAILABLE   AGE
deployment.apps/apache-deployment  0/1      1              0           102s

NAME                                DESIRED    CURRENT    READY   AGE
replicaset.apps/apache-deployment-6d9b8fb897  1          1          0       102s
vagrant@k8s-master:~$
```

Noticeable Changes :

- Pod created but container crashing , Reason dont have enough source to run apache





## Situation 3.2 : What if I set limit memory to 10Mi than required 6Mi and keep cpu same

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apache-deployment
  labels:
    app: apache
spec:
  replicas: 1
  selector:
    matchLabels:
      app: apache
  template:
    metadata:
      labels:
        app: apache
    spec:
      containers:
        - name: apache
          image: httpd:latest
          ports:
            - containerPort: 80
          resources:
            limits:
              memory: "10Mi"
              cpu: "10m"
            requests:
              memory: "3Mi"
              cpu: "1m"
```

```
vagrant@k8s-master:~$ kubectl get all -n vishalk17
NAME                                READY   STATUS    RESTARTS   AGE
pod/apache-deployment-6f6c4cccb8-pl7xn  1/1     Running   0           20s

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/apache-deployment  1/1     1             1           20s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/apache-deployment-6f6c4cccb8  1         1         1       20s
vagrant@k8s-master:~$
vagrant@k8s-master:~$ kubectl top pods -n vishalk17
NAME                                CPU(cores)   MEMORY(bytes)
apache-deployment-6f6c4cccb8-pl7xn  5m           6Mi
vagrant@k8s-master:~$
```

Noticeable Changes :

- Container created,

## Situation 3.2 : What if I set limit not a request

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: apache-deployment
  labels:
    app: apache
spec:
  replicas: 1
  selector:
    matchLabels:
      app: apache
  template:
    metadata:
      labels:
        app: apache
    spec:
      containers:
        - name: apache
          image: httpd:latest
          ports:
            - containerPort: 80
          resources:
            limits:
              memory: "10Mi"
              cpu: "10m"
```

```
vagrant@k8s-master:~$ kubectl apply -f httpd.yml -n vishalk17
deployment.apps/apache-deployment created
vagrant@k8s-master:~$ 
vagrant@k8s-master:~$ kubectl get all -n vishalk17
NAME                                READY   STATUS              RESTARTS   AGE
pod/apache-deployment-8dd5754dd-94qsz 0/1     ContainerCreating   0           9s

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/apache-deployment 0/1     1             0           9s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/apache-deployment-8dd5754dd 1         1         0       9s
vagrant@k8s-master:~$ kubectl get all -n vishalk17
NAME                                READY   STATUS    RESTARTS   AGE
pod/apache-deployment-8dd5754dd-94qsz 1/1     Running   0           15s

NAME                                READY   UP-TO-DATE   AVAILABLE   AGE
deployment.apps/apache-deployment 1/1     1             1           16s

NAME                                DESIRED   CURRENT   READY   AGE
replicaset.apps/apache-deployment-8dd5754dd 1         1         1       16s
vagrant@k8s-master:~$ kubectl top pods -n vishalk17
```

```
vagrant@k8s-master:~$ kubectl top pods -n vishalk17
NAME                                CPU(cores)   MEMORY(bytes)
apache-deployment-8dd5754dd-94qsz 1m           6Mi
vagrant@k8s-master:~$
```

Noticeable Changes :

- Container created, Limits = request

### Summary :

request = mentioned

limit = mentioned

no issues ( source can be used upto limits , above that container fail)

---

request = not mentioned

limit = mentioned

limit = request ( source can be used upto limits , above that container fail)

---

request = mentioned

limit = not mentioned

unlimited resources (default to zero) ( no limits defined, container will not fail and will not has limitation on using compute resources )

---

My devops repo :

- <https://github.com/vishalk17/devops>

My telegram channel:

-  [https://t.me/vishalk17\\_devops](https://t.me/vishalk17_devops)

Contact:

Telegram :  t.me/vishalk17

vishalk17 My youtube Channel :

-  **YouTube** <https://www.youtube.com/@vishalk17>